

We Reach AGI!!!

FirstName LASTNAME

Example University

Problem Formulation: A Curious Little Problem

The goal of this work is to understand something interesting but possibly unnecessary. But don't worry—we'll formalize it:

Definition (The Problem). Formally, this is the problem we're pretending is novel. It involves mysterious variables, vague assumptions, and a deep sense of academic urgency.

Here's a visual representation, just to look professional:




Figure 1. A figure. Definitely helps, right?

As you can see, this figure captures the essence of the problem.

Insightful Highlight Box. This box exists to draw your attention. What to? Not entirely sure. But it's highlighted, so it must be important.

Existing Approaches: A Brief Stroll Through Literature

A quick literature review to show we read a few papers:

- Classical physics seems somehow relevant [newton1726philosophiae].
- Honestly, we just needed a second bullet point.

Overview and Contributions

This work marks a pivotal moment in the history of ML and AI:

- A theoretical argument for why **AGI** is inevitable.
- We actually build such a model with *≈one trillion* parameters.
- Empirical results showing perfect scores on this year's IMO (not even out).

Theoretical Results: Unquestionable Mathematical Rigor

Now that we've stated the problem, let's prove some theorems. Or at least write some that look impressive:

Theorem (AGI Emergence). If your model has at least 10^{12} parameters and a sufficiently dramatic name, it will become sentient with probability $\epsilon > 0$.

Definition (Superintelligence Gradient). The partial derivative of model confidence with respect to Twitter hype. Empirically observed to be strictly increasing.




Figure 2. Evidence of emergent intelligence.

Failed Approaches: We Suffered So You Don't Have To

Naturally, we tried some things that didn't work:

- Copying ChatGPT:** Ethically dubious, technically tempting.
- Brute Force AGI:** Required more GPUs than the planet currently owns.

Theorem (Overfit Paradox). As model size increases, accuracy increases on everything except the test set.

Proposed Method: Definitely Not Just a Bigger Transformer

Our model is designed with scalability in mind.

Definition (Trillion-Parameter Transformer). A neural network so large, it requires its own power grid and a team of therapists.

Theorem (Universal Solver). Given enough data and a sufficiently vague benchmark, our model achieves state-of-the-art on at least one metric.

Experimental Results: Trust Us, We Have Graphs

We tested our model on:

- IMO. We got a perfect score. Somehow.
- The Turing Test. The judges asked it for dating advice and were convinced.
- StackOverflow. It answered "it depends" with 95% accuracy.




Figure 3. Blue line = us. Orange line = others. Need we say more?

Conclusions and Next Steps

- Our method is clearly on track to surpass human intelligence.
- Future work: deploy model on Mars and let it evolve in isolation.

Final Thought. This project started as a joke. Now it's under review at NeurIPS.